

Lab 4: Deep 3D Descriptors

Task 1: Anchors Generation

In the first task it is required to generate an anchor and a positive and a negative sample. To do this, a random point is extracted from the starting point cloud, together with its neighborhood using the `Open3d search_radius_vector_3d()` function, specifying the *radius*. This operation is performed also for the positive and negative sample, using the noisy point cloud. With the `search_knn_vector_3d()` function, the closest point to the previous one is extracted from the noisy point cloud, with its neighborhood. Regarding the negative sample, a random point from the noisy point cloud is selected, making sure it is at a certain distance from the anchor. This check is performed in a simple while loop. Finally the samples are normalized with respect to the center of the sphere. The figure below shows an example of the 3 samples.




Figure 1: Anchor, positive and negative samples

Task 2: TinyPointNet Structure

The structure of Tiny PointNet is represented in Figure 2. In order to construct the network, the already defined MLP and TNet classes are used in the `__init__()` method, along with the `forward()` method in the `TinyPointNet` class. The first layer employed is a TNet with input $k = 3$, followed by two shared MLPs that map input size 3 to 64 and then 64 to 64. Then another TNet is used, this time with $k = 64$ and again some MLPs that maps input size 64 to 256 with an intermediate dimension of 128. At the end a max pooling operation is performed to generate the final global features, a vector of dimension 256.

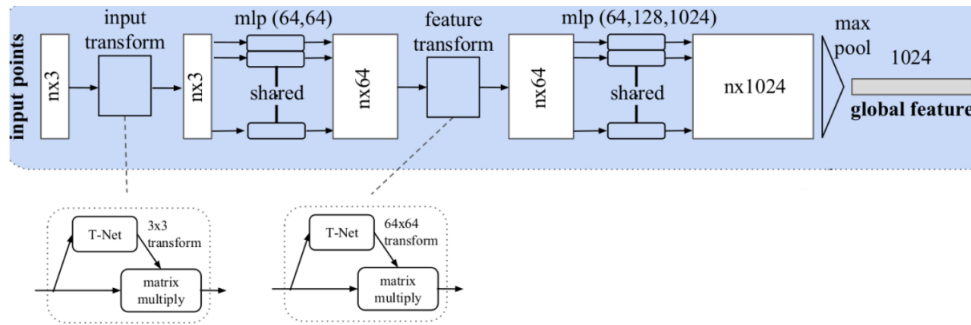


Figure 2: TinyPointNet structure

Task 3: Loss Function

The training operation is performed with a Triplet Loss Function with default parameters as margin equals to 1.0 (α in the equation) and p , the norm degree for pairwise distance as 2. In Pytorch, the loss is defined as `nn.TripletMarginLoss()`.

$$L(A, P, N) = \max(\|f(A) - f(P)\| - \|f(A) - f(N)\| + \alpha, 0) \quad (1)$$

Results

Many test have been carried out, trying different values of learning rate and radius, in order to obtain good results and prevent over-fitting. The best values seems to be learning rate equals to 0.001 and radius $3 \cdot 10^{-3}$. This two hyper-parameters allowed the network to reach a lower loss value, both on train and validation sets, as shown in Figure 3.

Regarding the test part, the network is able to achieve an accuracy value of 88.235% on the test set. However, it is important to highlight that different runs produced different accuracy values, most of them in the range 80 – 90%, and only once it performed a lower value of 50%.




Figure 3: Train and validation losses